

Credit Risk Scoring

- Objective 2.1: Data Cleaning
- Objective 2.2: Feature Scaling & Hyper-Parameter Tuning
- Objective 2.3 Score-carding *with* Machine Learning

Problem and Motivation

- **Environmental Scope:**
 - One of the Company's services ("Loan Plan" Service), provided to new eligible customers seeking a monetary loan.
- **Problems:**
 - Before Granting such Loans, the Company would First Need to Evaluate All new Applicants based on their Submitted Loan Registration Details:
 - Loan History
 - Assets Currently Available (e.g. Owned/Rented Properties)
 - and more
 - These details allow the Company to Tabulate the Probability of these applicants defaulting on their Desired Loans through referencing the Company's Historical Loan Lending Data of Previous Customers to date with their Loan Outcomes (i.e. whether they have Defaulted or Not).
 - From there, the Company can then Grant Customers displaying Lower Probabilities of Defaulting, on their Requested Loans, while Considering the Company's Currently Available Funding for provisioning the New Loans.
 - This is known as **Credit Scoring**.

Problem and Motivation

- **Motivation:**
 - Digitalizing and Automating an efficient Credit Scoring System
 - To reduce any unnecessary Operating Expenditure(s) caused by Human intervention (i.e. Credit Scoring Mistakes, Slow Score Tabulations, and other Human-Related Errors).
 - Despite it being possible to curate a Algorithmic System (Non-Machine Learning) for such a purpose, the Company has personally Requested the Use of Machine Learning, particularly:
 - Logistic Regression (LR) algorithm
 - With a Measurable Scorecard for more Transparent Credit Scoring Interpretations.

GAPS Capstone Tackles

Studies Lack of:

- **Feature Scaling** (*Gap 1*)
 - Nikolic et al [15] : **No Scaling indications** in Study
 - Li and Chen [12] : **No Scaling indications** in Study
 - Bensic et al [13] : **No Scaling indications** in Study
- **Hyper-Parameter Tuning of Logistic Regression Model** (*Gap 2*)
 - Bensic et al [13] : Tuned all Models **Except** for **Logistic Regression (LR)** Using Credit Scoring Data, experimented using **LR**, 4 Neural Networks, Decision Trees. Study's Justification for **Not** tuning **LR- Small Dataset**
 - Li and Chen [12] : Tuning Done on Random Forest, except for **LR (No Tuning Indications** in Study)
(Aim of Study- Use 10 Models Classifiers, Evaluated using AUC/Accuracy/Kolmogorov-Smirnov Performance Metrics)

As such, the Capstone's solution would include the use of both Feature Scaling and Hyper-Parameter Tuning to Improve Credit Scoring Performance in Continuation to the Studies.

Objective 2.1

Clean and Transform the Dataset

Raw (Loan) Data

- Dataset from **Lending Club**TM (peer-to-peer Lending Company based in the US).
 - Loans Data from Year 2007 to 2018 (4th Quarter).
 - 2 Million Observations (Rows) and 145 Features (Variables).

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc
0	NaN	NaN	2500	2500	2500.0	36 months	13.56	84.92	C	C1	Chef	10+ years	RENT	5
1	NaN	NaN	30000	30000	30000.0	60 months	18.94	777.23	D	D2	Postmaster	10+ years	MORTGAGE	9
2	NaN	NaN	5000	5000	5000.0	36 months	17.97	180.69	D	D1	Administrative	6 years	MORTGAGE	5
3	NaN	NaN	4000	4000	4000.0	36 months	18.94	146.51	D	D2	IT Supervisor	10+ years	MORTGAGE	9
4	NaN	NaN	30000	30000	30000.0	60 months	16.14	731.78	C	C4	Mechanic	10+ years	MORTGAGE	5

Columns Kept (after Dropping 80% Null Columns)

**Columns Kept
(97 columns):**

#	Column	Dtype	24	revol_util	float64	45	open_il_24m	float64	66	mo_sin_rcnt_tl	float64	86	pct_tl_nvr_dlq	float64
---	-----	-----												
0	loan_amnt	int64	25	total_acc	float64	46	mths_since_rcnt_il	float64	67	mort_acc	float64	87	percent_bc_gt_75	float64
1	funded_amnt	int64	26	initial_list_status	object	47	total_bal_il	float64	68	mths_since_recent_bc	float64	88	pub_rec_bankruptcies	float64
2	funded_amnt_inv	float64	27	out_prncp	float64	48	il_util	float64	69	mths_since_recent_bc_dlq	float64	89	tax_liens	float64
3	term	object	28	out_prncp_inv	float64	49	open_rv_12m	float64	70	mths_since_recent_inq	float64	90	tot_hi_cred_lim	float64
4	int_rate	float64	29	total_pymnt	float64	50	open_rv_24m	float64	71	mths_since_recent_revol_delinq	float64	91	total_bal_ex_mort	float64
5	installment	float64	30	total_pymnt_inv	float64	51	max_bal_bc	float64	72	num_accts_ever_120_pd	float64	92	total_bc_limit	float64
6	grade	object	31	total_rec_int	float64	52	all_util	float64	73	num_actv_bc_tl	float64	93	total_il_high_credit_limit	float64
7	emp_length	object	32	last_pymnt_d	object	53	total_rev_hi_lim	float64	74	num_actv_rev_tl	float64	94	hardship_flag	object
8	home_ownership	object	33	last_pymnt_amnt	float64	54	inq_fi	float64	75	num_bc_sats	float64	95	disbursement_method	object
9	annual_inc	float64	34	last_credit_pull_d	object	55	total_cu_tl	float64	76	num_bc_tl	float64	96	debt_settlement_flag	object
10	verification_status	object	35	collections_12_mths_ex_med	float64	56	inq_last_12m	float64	77	num_il_tl	float64			
11	issue_d	object	36	mths_since_last_major_derog	float64	57	acc_open_past_24mths	float64	78	num_op_rev_tl	float64			
12	loan_status	object	37	policy_code	int64	58	avg_cur_bal	float64	79	num_rev_accts	float64			
13	pymnt_plan	object	38	application_type	object	59	bc_open_to_buy	float64	80	num_rev_tl_bal_gt_0	float64			
14	purpose	object	39	acc_now_delinq	float64	60	bc_util	float64	81	num_sats	float64			
15	addr_state	object	40	tot_coll_amt	float64	61	chargeoff_within_12_mths	float64	82	num_tl_120dpd_2m	float64			
16	dti	float64	41	tot_cur_bal	float64	62	delinq_amnt	float64	83	num_tl_30dpd	float64			
17	delinq_2yrs	float64	42	open_acc_6m	float64	63	mo_sin_old_il_acct	float64	84	num_tl_90g_dpd_24m	float64			

Utilize Stratification Split (Reduce Imbalance Classes)

Split data

```
# Split the data using an 80/20 split
X = loan_data.drop('good_bad', axis = 1)
y = loan_data['good_bad']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42, stratify = y)
X_train, X_test = X_train.copy(), X_test.copy()
```

- **Training:Test Split** (80:20) to **Prevent Overfitting**
(as compared to using 100% Dataset for Testing)
- **Random Stratified Sampling** also Assists in **Reducing Class imbalances** during Data Sampling.

Vannoy et al [75]

Wall and
Fontenot [77]

Data cleaning (Data-Specific)

- String Replacement (emp_length, **String to Numbers** only)

```
array([ 7., 10.,  3.,  4.,  2.,  0.,  1.,  6.,  5.,  8.,  9.])
```

E.g. emp_length String to Integers type

- **Date Conversions to Integer** (Number of Days from specific Date- e.g. Today)
(earliest_cr_line, issue_d, last_pymnt_d, last_credit_pull_d)

Feature Selection

- A p-value **less than 0.05** (typically ≤ 0.05) **is statistically significant.**

Panesar [53]

- It indicates Strong Evidence against the Null Hypothesis (No **Statistical Relationship**).

- P value **greater than 0.05** means (**Insignificant**) that no effect was observed.

Cai et al [21],
Panesar [53],
Doan et al [55]

- Remove such Insignificant Features (≥ 0.05)

Chi-Squared

- Gauges the Degree of Independency between Categorical Variables.
 - i.e. Identifies Highly Related/Dependent Variables (against One Another)
- Also gives the Option of Determining which Categories are Significant.
 - Being the “most useful tools in the researcher’s array of available analysis tools”
- Limited to **only** Discrete/Categorical Variables.
- **Action: Select top 5**

	Feature	p-value
0	grade	0.0
1	home_ownership	0.0
2	verification_status	0.0
3	pymnt_plan	0.0
4	purpose	0.0
5	addr_state	0.0
6	initial_list_status	0.0
7	application_type	0.0
8	hardship_flag	0.0
9	disbursement_method	0.0
10	debt_settlement_flag	0.0

Kamalov and Thabtah [44]

McHugh [78]

Urbanowicz [54]

ANOVA F-Score

- Determines Strength of Relationship between Features and Labels
- Able to handle Discrete and Continuous Variables.
 - Replaces Chi-Squared for Continuous Variables.
- **Action: Select top 20**

	Numerical_Feature	F-Score	p values
0	mths_since_last_pymnt_d	101279.076110	0.0
1	int_rate	79891.093938	0.0
2	mths_since_last_credit_pull_d	64883.628239	0.0
3	last_pymnt_amnt	63017.806787	0.0
4	out_prncp	60841.859961	0.0
5	out_prncp_inv	60830.025704	0.0
6	mths_since_issue_d	36090.598122	0.0
7	total_pymnt_inv	34518.813024	0.0
8	total_pymnt	34484.620171	0.0
9	acc_open_past_24mths	15827.049134	0.0
10	term	14166.382950	0.0
11	inq_last_6mths	13784.078774	0.0
12	num_tl_op_past_12m	12395.693952	0.0
13	bc_open_to_buy	11932.905087	0.0
14	total_bc_limit	10038.049778	0.0
15	percent_bc_gt_75	9830.099516	0.0
16	bc_util	9186.155895	0.0
17	revol_util	7627.581526	0.0
18	tot_hi_cred_lim	7152.718716	0.0
19	mths_since_recent_inq	6168.541817	0.0

Yange et al [79]

Urbanowicz
et al [54]

One Hot Encoding

- One-Hot Encoding was used as opposed to Label Encoding.
 - Label Encoding (uses numerical values eg. 5 , 7) to represent category labels
 - ML algorithms may perceive such labels as a form of numerical relation (e.g. hierarchy/order), leading potentially bias predictions.
 - One-Hot Encoding (uses binary values e.g. 1, 0) and separate columns instead
 - Eliminates undesirable numerical relation biases

grade:A	grade:B	grade:C	grade:D	grade:E	grade:F	grade:G	home_ownership:ANY	home_ownership:MORTGAGE	home_ownership:NONE	home_ownership:OTHE
0	0	1	0	0	0	0	0	1	0	
0	0	1	0	0	0	0	0	1	0	
0	0	0	1	0	0	0	0	0	0	
0	0	1	0	0	0	0	0	1	0	
0	0	1	0	0	0	0	0	0	0	

Types of Features

- Discrete (**Categorical**) Features ->

grade	home_ownership
C	MORTGAGE
C	MORTGAGE
D	RENT
C	MORTGAGE
C	RENT

- Continuous (**Numerical**) Features ->

inq_last_6mths	revol_util
2.0	36.4
0.0	90.7
0.0	44.0
0.0	39.3
0.0	39.0

Weight of Evidence (WOE)

- Estimates the **Predictive Power** (I.e. Weight) of any **Given Bin**, within a Feature.
 - Later Transformed to **Determine Information Value (IV)** - Main Goal.

Siddiqi [57] ,
Zeng [59]

$$WoE = \ln(\text{goods}_i / \sum_{i=1}^n \text{goods}_i) - \ln(\text{bads}_i / \sum_{i=1}^n \text{bads}_i)$$

Nikolic et
al's study [3]

Distribution of Non-Defaulting Cases DIVIDED BY Defaulting Cases

- **Inner Workings of WOE:**
 - Contribution of Each Feature in Relation to the Actual Outcome (Predictive Power)
 - Rank-ability of Features' Bins according to their Predictive Power
 - Accept and Handle Missing Values without the Need of any Imputative Actions, on the Dataset.

Siddiqi [57] and
Baesens et al's [58]

Information Value (IV)

- Determines the Feature's uncertainty of predicting the outcome based on its

Zhang et al
[81]

Sum of the (% of non-defaulting customers - % of defaulting customers)

*

WOE

(WOE)



$$IV = \sum_{i=1}^n (\text{Distr Good}_i - \text{Distr Bad}_i) * \ln \left(\frac{\text{Distr Good}_i}{\text{Distr Bad}_i} \right)$$

Nikolic et al [15]

Determining Information Value

Guideline for Information Value (IV):

- $IV < 0.02$ == not useful for prediction (Action: Ignore Feature)
- $IV 0.02$ to 0.1 == weak predictive power **(Action: Keep Feature)**
- $IV 0.1$ to 0.3 == medium predictive power **(Action: Keep Feature)**
- $IV 0.3$ to 0.5 == strong predictive power **(Action: Keep Feature)**
- $IV > 0.5$ suspicious predictive power (Action: Ignore Feature)

WOE (Pseudo Code 1)

- Pseudo-Code of **Number of Observations** [*'n_obs' column*] **AND** **Proportion of Non-Defaulting** [*'prop_good' column*]
- For Each Feature (eg. 'Grade' Feature):
 - For Each Feature's Unique Value/Value-Range (eg. Categorical Value of 'A'/'B'/'C'/'D' **OR** Continuous Value-Range of (5.284, 6.594] / (6.594, 7.878])
 - Calculate the **Number of Observations** (i.e. '*n_obs*' column)
Get the **Number of Non-Defaulting Records** that Uses that Feature's Value/Value-Range (eg. Total of 12,000 Defaulting Cases)
 - Calculate the **Proportion of Non-Defaulting** (i.e. '*prop_good*' column)
Get the **Proportion of the Number of Non-Defaulting Records** that uses that Feature's Categorical/Continuous Value/Value-Range
(eg. Proportion of 0.6 for Grade 'B' means that 60% of All Records that uses the Feature's Categorical/Continuous Value/Value-Range of Grade 'B' are non-Defaulting Outcomes)

WOE (Pseudo Code 2)

- Pseudo-Code of **WOE** (i.e. '**WoE**' column in [Section 3.6](#)):
- For Each Feature (eg. 'Grade' Feature):
 - For Each Feature's Unique Value/Value-Range (eg. Categorical Value of 'A'/'B'/'C'/'D' **OR** Continuous Value-Range of (5.284, 6.594] / (6.594, 7.878])
 - Calculate **Number of Observations**
 - Calculate **Proportion of Non-Defaulting Records**
 - Calculate **Proportion of Defaulting Records** (i.e. Same Procedure as **Proportion of Non-Defaulting Records** but by using Defaulting Outcomes)
 - Calculate the **Number of Non-Defaulting Records**
 - (i.e. '**n_good**' column: **Proportion of Non-Defaulting Records** MULTIPLY **Number of Observations**)
 - Calculate the **Number of Defaulting Records**
 - (i.e. '**n_bad**' column: [1 MINUS **Proportion of Defaulting Records**] MULTIPLY **Number of Observations**)
 - Calculate the **Proportion of Non-Defaulting Records**
 - (i.e. '**prop_n_good**' column: **Number of Non-Defaulting Records** DIVIDED BY **Summation of All Non-Defaulting Records in Feature**)
 - Calculate the **Proportion of Defaulting Records**
 - (i.e. '**prop_n_bad**' column: **Number of Defaulting Records** DIVIDED BY **Summation of All Defaulting Records in Feature**)
 - Calculate **WOE** of Each Feature's Categorical/Continuous Value/Value-Range
(**Proportion of Non-Defaulting Records** DIVIDED BY **Proportion of Defaulting Records**)

Example 1: Information Value

- $IV < 0.02$ == not useful for prediction (Action: Ignore Feature)
- IV 0.02 to 0.1 == weak predictive power (**Action: Keep Feature**)
- IV 0.1 to 0.3 == medium predictive power (**Action: Keep Feature**)
- IV 0.3 to 0.5 == strong predictive power (**Action: Keep Feature**)
- $IV > 0.5$ suspicious predictive power (Action: Ignore Feature)

Out[]:

	purpose	n_obs	prop_good	prop_n_obs	n_good	n_bad	prop_n_good	prop_n_bad	WoE	diff_prop_good	diff_WoE	IV
0	small_business	5605	0.787333	0.015026	4413.0	1192.0	0.013282	0.029234	-0.788933	NaN	NaN	0.036668
1	educational	351	0.792023	0.000941	278.0	73.0	0.000837	0.001790	-0.760693	0.004690	0.028240	0.036668
2	renewable_energy	295	0.837288	0.000791	247.0	48.0	0.000743	0.001177	-0.459668	0.045265	0.301026	0.036668
3	moving	2397	0.848561	0.006426	2034.0	363.0	0.006122	0.008903	-0.374498	0.011273	0.085169	0.036668
4	house	1824	0.861294	0.004890	1571.0	253.0	0.004728	0.006205	-0.271777	0.012733	0.102721	0.036668
5	other	19006	0.861675	0.050951	16377.0	2629.0	0.049291	0.064477	-0.268581	0.000381	0.003196	0.036668
6	medical	3750	0.863467	0.010053	3238.0	512.0	0.009746	0.012557	-0.253469	0.001791	0.015112	0.036668

Example 2: Information Value

- $IV < 0.02$ == not useful for prediction (**Action: Ignore Feature**)
- IV 0.02 to 0.1 == weak predictive power (**Action: Keep Feature**)
- IV 0.1 to 0.3 == medium predictive power (**Action: Keep Feature**)
- IV 0.3 to 0.5 == strong predictive power (**Action: Keep Feature**)
- $IV > 0.5$ suspicious predictive power (Action: Ignore Feature)

Out[]:	emp_length	n_obs	prop_good	prop_n_obs	n_good	n_bad	prop_n_good	prop_n_bad	WoE	diff_prop_good	diff_WoE	IV
0	0.0	45764	0.876497	0.122682	40112.0	5652.0	0.120727	0.138618	-0.138189	NaN	NaN	0.006408
1	1.0	23582	0.888262	0.063218	20947.0	2635.0	0.063045	0.064625	-0.024743	0.011765	0.113446	0.006408
2	2.0	33123	0.889895	0.088795	29476.0	3647.0	0.088715	0.089444	-0.008184	0.001633	0.016559	0.006408
3	3.0	29301	0.890925	0.078549	26105.0	3196.0	0.078569	0.078383	0.002372	0.001030	0.010555	0.006408
4	4.0	22482	0.891602	0.060269	20045.0	2437.0	0.060330	0.059768	0.009357	0.000677	0.006985	0.006408
5	5.0	24654	0.885374	0.066092	21828.0	2826.0	0.065697	0.069309	-0.053524	0.006229	0.062881	0.006408
6	6.0	21057	0.883079	0.056449	18595.0	2462.0	0.055966	0.060382	-0.075936	0.002294	0.022413	0.006408

Steps for WOE Binning

Only if the Feature's **IV is within the Acceptable Range**, then it'll be allowed to be WOE Binned:

- Each bin should contain 5% of the observations (*prop_n_obs*)
 - Fewer the bins increases 'smoothing' allowing the capture of important data while leaving out noise.
- The Current Approach taken: 10/20 bins
 - Too Few Bins, Reduce Data Distribution leading to Model Instability)
 - Too Many Bins, causes Model Overfitting
- Each Category (Bin) must be Non-Zero for both events and non-events (e.g. default and non-default)
- WOE Values Must be Distinct for each Category
 - Similar Groups are Binned Together (I.e. Bins with Similar WOE values have Same Proportion of Defaulting/Not-Defaulting on loans -> Same predictive power)
- Missing Values are Binned Separately

Eg: Plan Weight Of Evidence (**Continuous** Features)

- Each Bin
5% of Observations

	int_rate_factor	n_obs		prop_n_obs	n_good	n_bad	prop_n_good	prop_n_bad	WoE	diff_prop_good	diff_WoE	IV
0	(5.284, 6.594]	99983	Bin 1									
			Bin 2	0.055284	97848.0	2135.0	0.061888	0.009386	1.886139	NaN	NaN	0.449621
1	(6.594, 7.878]	145963	Bin 3	0.080708	141215.0	4748.0	0.089317	0.020873	1.453750	0.011175	0.432389	0.449621
2	(7.878, 9.162]	154733	Bin 4	0.085557	144831.0	9902.0	0.091604	0.043530	0.744021	0.031465	0.709729	0.449621
3	(9.162, 10.446]	166225	Bin 5	0.091911	154837.0	11388.0	0.097932	0.050063	0.671003	0.004515	0.073018	0.449621
4	(10.446, 11.73]	217977	Bin 6	0.120527	197517.0	20460.0	0.124927	0.089944	0.328543	0.025354	0.342460	0.449621
5	(11.73, 13.014]	201025	Bin 7	0.111154	176975.0	24050.0	0.111934	0.105726	0.057064	0.025774	0.271479	0.449621
6	(13.014, 14.298]	187860	Bin 8	0.103874	161992.0	25868.0	0.102458	0.113718	-0.104270	0.018061	0.161333	0.449621
7	(14.298, 15.582]	142739	Bin 9	0.078925	121497.0	21242.0	0.076845	0.093382	-0.194901	0.011119	0.090631	0.449621
8	(15.582, 16.866]	118746	Bin 10	0.065659	97265.0	21481.0	0.061519	0.094432	-0.428539	0.032082	0.233639	0.449621
9	(16.866, 18.15]	113345	Bin 11	0.062672	91132.0	22213.0	0.057640	0.097650	-0.527179	0.015078	0.098639	0.449621
10	(18.15, 19.434]	79641	Bin 12	0.044036	61614.0	18027.0	0.038970	0.079248	-0.709791	0.030376	0.182613	0.449621
11	(19.434, 20.718]	48931		0.027056	37599.0	11332.0	0.023781	0.049816	-0.739463	0.005238	0.029672	0.449621
12	(20.718, 22.002]	39421		0.021797	29906.0	9515.0	0.018915	0.041829	-0.793620	0.009777	0.054157	0.449621
13	(22.002, 23.286]	21935	Bin 13	0.012129	16163.0	5772.0	0.010223	0.025374	-0.909104	0.021772	0.115484	0.449621
14	(23.286, 24.57]	21226		0.011737	15596.0	5630.0	0.009864	0.024750	-0.919905	0.002100	0.010801	0.449621
15	(24.57, 25.854]	20624		0.011404	14780.0	5844.0	0.009348	0.025691	-1.010950	0.018118	0.091046	0.449621
16	(25.854, 27.138]	10973		0.006067	7950.0	3023.0	0.005028	0.013289	-0.971888	0.007865	0.039063	0.449621
17	(27.138, 28.422]	4060		0.002245	3060.0	1000.0	0.001935	0.004396	-0.820395	0.029189	0.151493	0.449621
18	(28.422, 29.706]	5406		0.002989	3896.0	1510.0	0.002464	0.006638	-0.990969	0.033014	0.170574	0.449621
19	(29.706, 30.99]	7721		0.004269	5386.0	2335.0	0.003407	0.010265	-1.103019	0.023103	0.112050	0.449621

Eg: Plan Weight Of Evidence (**Discrete** Features)

- Each Bin 5% of Observations

	grade	n_obs	prop_good	prop_n_obs	n_good	n_bad	prop_n_good	prop_n_bad	WoE	diff_prop_good	diff_WoE	IV
0	G	9697	Bin 1	0.005362	5845.0	3852.0	0.003697	0.016934	-1.521816	NaN	NaN	0.483175
1	F	33507		0.018527	21455.0	12052.0	0.013570	0.052982	-1.362083	0.037550	0.159733	0.483175
2	E	108512		0.060000	78263.0	30249.0	0.049500	0.132977	-0.988198	0.080924	0.373885	0.483175
3	D	259341	Bin 2	0.143398	207569.0	51772.0	0.131285	0.227594	-0.550195	0.079133	0.438003	0.483175
4	C	519664	Bin 3	0.287340	446903.0	72761.0	0.282661	0.319864	-0.123648	0.059614	0.426547	0.483175
5	B	531178	Bin 4	0.293706	486447.0	44731.0	0.307672	0.196641	0.447651	0.055805	0.571300	0.483175
6	A	346635	Bin 5	0.191666	334577.0	12058.0	0.211616	0.053008	1.384329	0.049425	0.936678	0.483175

Objective 2.2

Feature Scaling & Hyper-Parameter Tuning
(Tackle Gaps here)

Machine Learning Pipelines (under Evaluation)

4 Pipelines:

- Pipeline 1: WOE Binning > Logistic Regression (**Untuned**)
- Pipeline 2: WOE Binning > Logistic Regression (**Tuned**)
- Pipeline 3: WOE Binning > **StandardScaler** > Logistic Regression (**Untuned**)
- Pipeline 4: WOE Binning > **StandardScaler** > Logistic Regression (**Tuned**)

Aim of Pipelines: Tackle (two) Gaps in Studies

*(**Tuned**) means that Logistic Regression Model being Trained by GridSearch Hyper-Parameter Optimization

** **StandardScaler** means that Dataset is Normalized through Feature Scaling

Logistic Regression (Tuning and Parameters)

- **Brute Force approach** (discovery of 14 million Models) by W. Książek et al's study [11]
 - **GridSearch** Hyper-Parameter Optimization used (for Experiment)
- A Few **Parameters adapted** from W. Książek et al's study [11]:
 - **C**: 1-100
("Higher values generalize the model")
 - **Max Iteration** : 1000
("Maximum iterations to converge the model")
 - **Penalty**: l2
"Adds bias to Model when it is suffering from high variance"
 - **Solver**: lbfgs, liblinear, saga
("Type of an Algorithm Solver")

Pipeline (Model Evaluation)

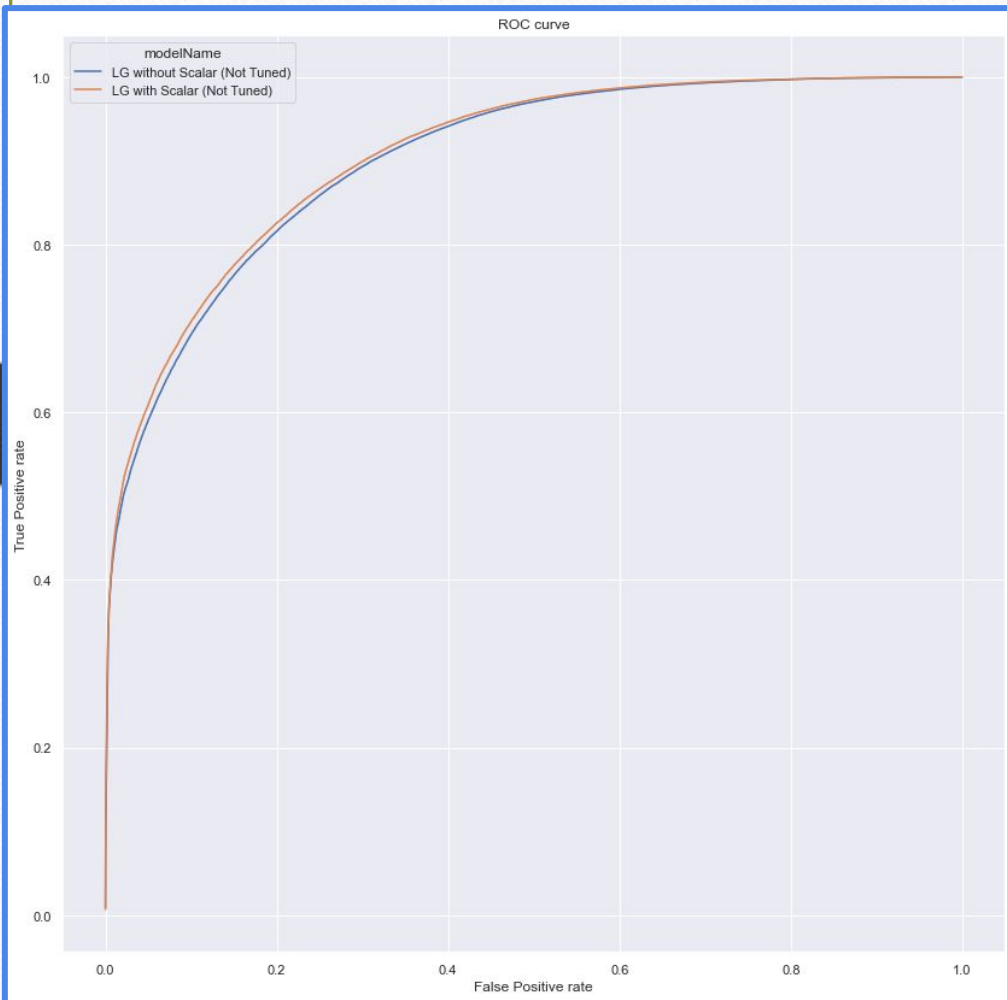
Evaluate Models via **Performance Metrics**:

- Reason: Attain **Numerical** TP, TN, FP, FN.
- **ROC Curve** (Graphical Plot): *Relation between TPR/FPR*
 - Reason: **Visually Compare** 4 Pipelines' TPR and FPR.
- **AUROC** (Performance Score): *Area Under ROC Curve*, closer to 1 means better performance
 - Reason: Should TPR/FPR **Results be Too Small** for ROC Curve Inspection, a Score of 0 to 1 is Given for Easier Referencing (Based on TPR/FPR Values).
- **GINI** (Performance Score): *Twice of Area Under ROC*
 - **Reaffirm** AUROC score (Based on TPR/FPR Values).

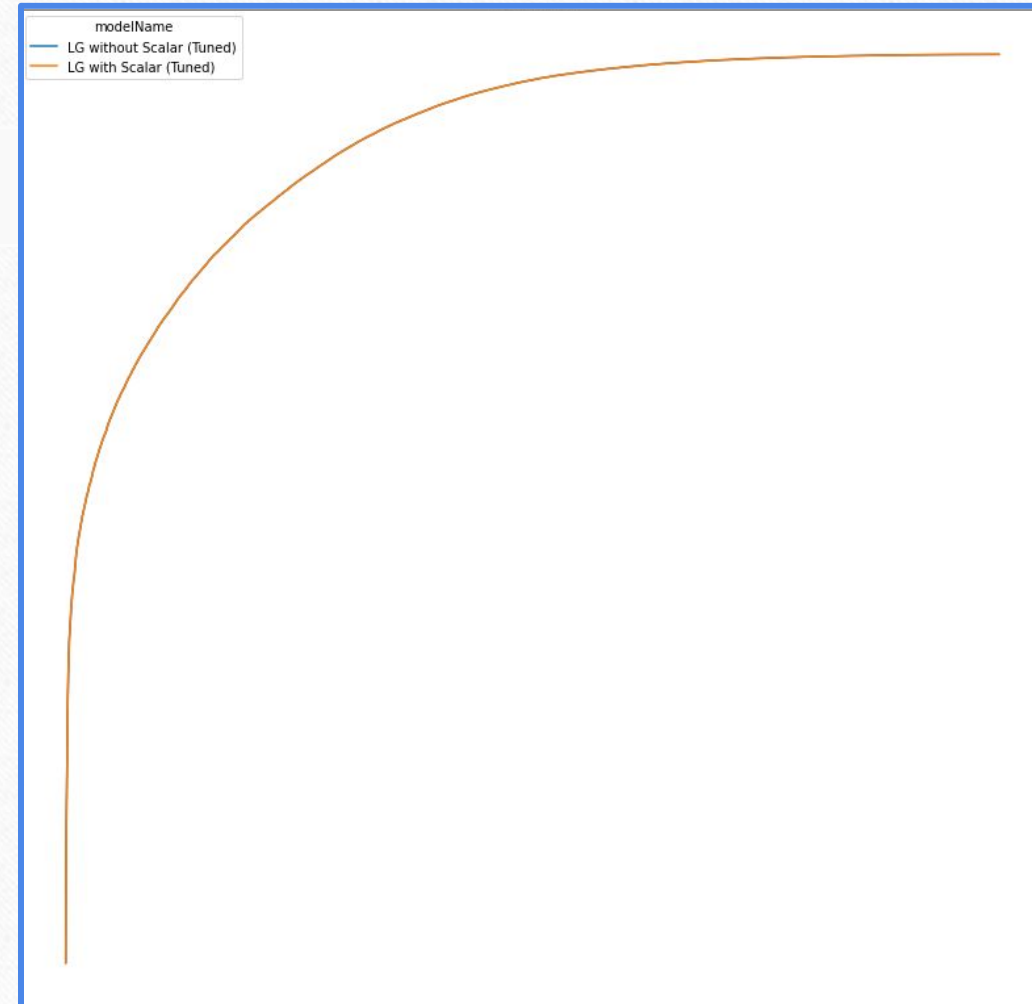
Results

Pipeline No.	1	2	3	4
LR Tuned?	No	Yes	No	Yes
Feature Scaling?	No	No	Yes	Yes
GINI Score	0.8158713758013114	0.8158900854661821	0.8158908685747437	0.8158908685747437
AUROC Score	0.9079356879006557	0.9079450427330911	0.9079454342873718	0.9079454342873718
True Positive (TP)	26107	26121	26120	26122
True Negative (TN)	387327	387327	387328	387328
False Positive (FP)	30762	30748	30749	30747
False Negative (FN)	7938	7938	7937	7937
Confusion Matrix Original Plot Found in	Figure 6 (Appendix A)	Figure 7 (Appendix A)	Figure 6 (Appendix A)	Figure 7 (Appendix A)
Confusion Matrix Plot Name (in Figure)	LG (No Scalar)	LG (No Scalar) Tuned	LG & Scalar	LG & Scalar Tuned
ROC Original Plot Found in	Figure 29 <u>LG without Scalar (Not Tuned)</u>	Figure 30 <u>LG without Scalar (Tuned)</u>	Figure 29 <u>LG with Scalar (Not Tuned)</u>	Figure 30 <u>LG with Scalar (Tuned)</u>

ROC Curves (Visual but too small Value Variances)

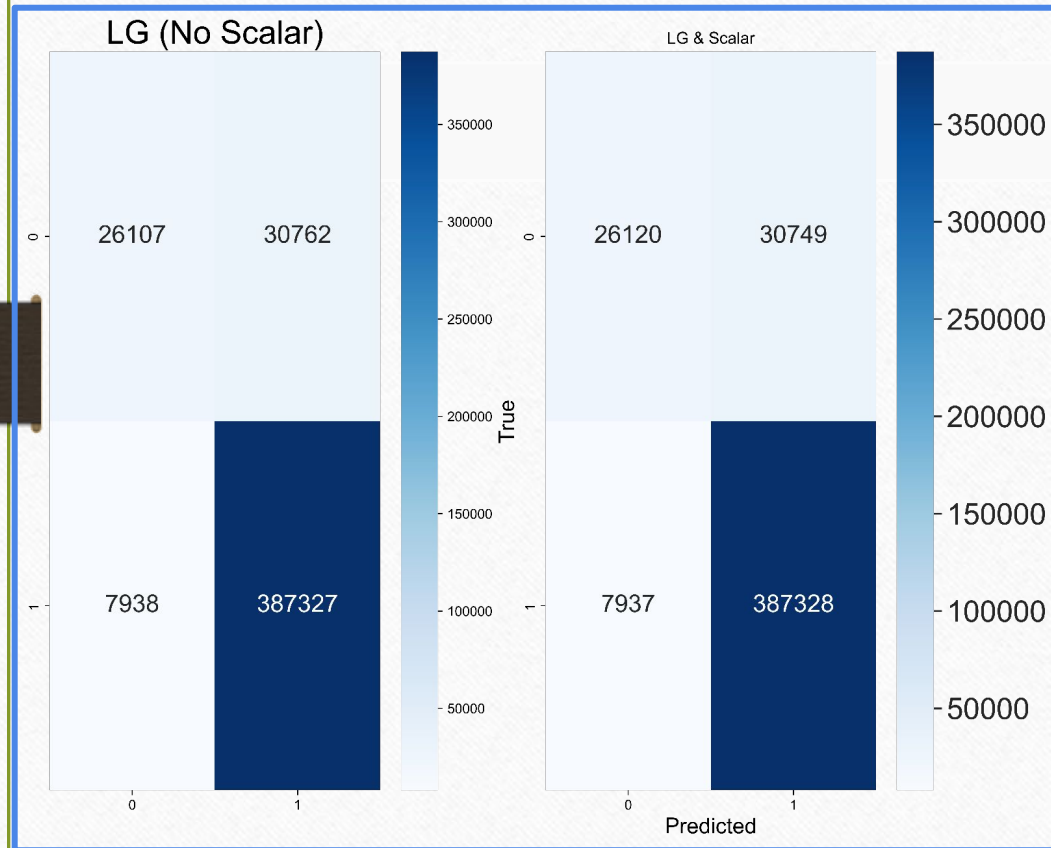


ROC Curve for **Pipeline 1 and 3**, Without and With Feature Scaling Respectively

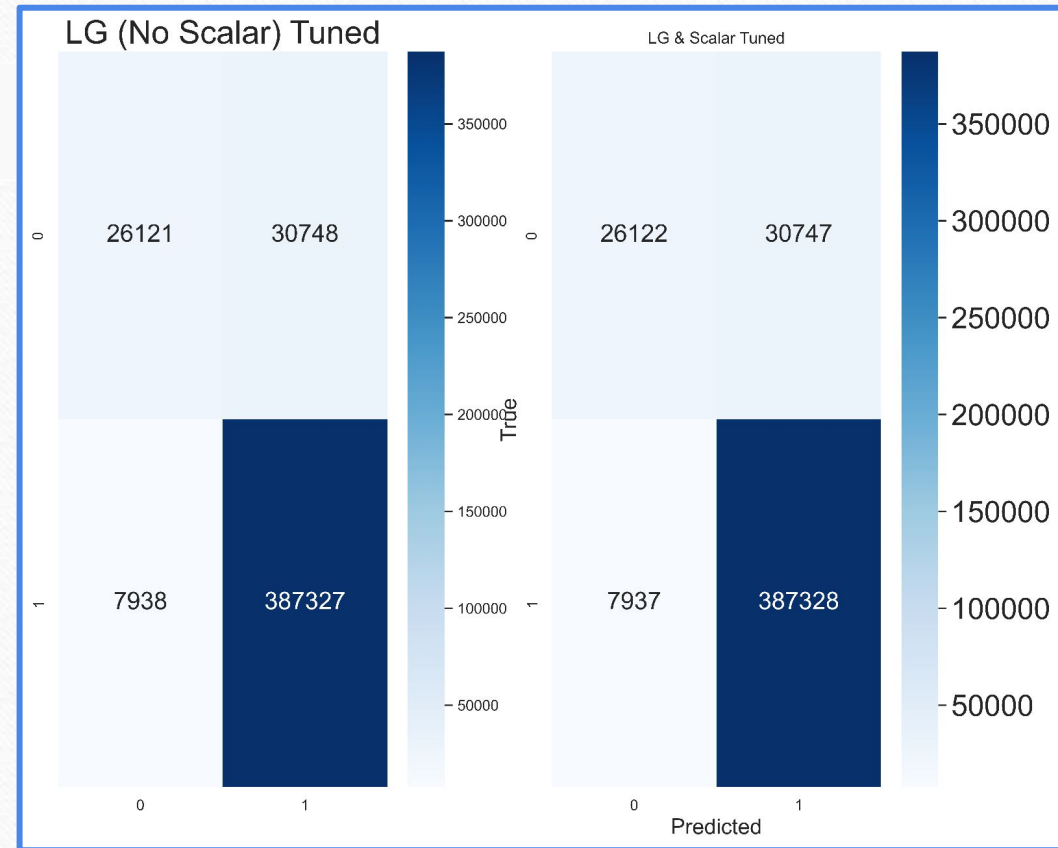


ROC Curve for **Pipeline 2 and 4**, Without and With Feature Scaling Respectively

ROC Curves (Visual but too small Value Variances)



Confusion Matrixes for **Untuned** Pipeline 1 and 3,
Without and With Feature Scaling Respectively



Confusion Matrixes for **Tuned** Pipeline 2 and 4,
Without and With Feature Scaling Respectively

Discussion: Feature Scaling (GAP)

Comparing Both **Untuned** LR Models (Pipeline 1 and 3):

- Pipeline 3 (Feature **Scaled** and **Without Tuning**) had **13 TP and 1 TN** (how the Model “got it right”) Values more than Pipeline 1 (**No** Feature Scaling and **Without Tuning**).
- Pipeline 3 had **13 FP and 1 FN** (how the Model “got it wrong”) Values lesser than Pipeline 1.
- Pipeline 3 has a **Higher** AUROC Score when compared to Pipeline 1.
- Pipeline 3 has a **Higher** GINI Score when compared to Pipeline 1.

Comparing Both **Tuned** LR Models (Pipeline 2 and 4):

- Pipeline 4 (Feature **Scaled** and **With Tuning**) also had **1 TP and 1 TN** Values more than the Pipeline 2 (**No** Feature Scaling and **With Tuning**).
- Pipeline 4 had **1 FP and 1 FN** Values lesser than Pipeline 2.
- Pipeline 4 has a **Higher** AUROC Score when compared to Pipeline 2.
- Pipeline 4 has a **Higher** GINI Score when compared to Pipeline 2.

Discussion: Hyper-Parameter Tuning (GAP)

Comparing Both (Feature) **Unscaled** LR Models (Pipeline 1 and 2):

- Pipeline 2 (**Tuning** and **Without Feature Scaling**) had **14 TP** (how the Model “got it right”) Values more than Pipeline 1 (**No** Tuning and **Without Feature Scaling**).
- Pipeline 2 had **14 FP** (how the Model “got it wrong”) Values lesser than Pipeline 1.
- Pipeline 2 has a **Higher** AUROC Score when compared to Pipeline 1.
- Pipeline 2 has a **Higher** GINI Score when compared to Pipeline 1.

Comparing Both (Feature) **Scaled** LR Models (Pipeline 3 and 4):

- Pipeline 4 (Feature **Scaled** and **With Tuning**) also had **2 TP** values more than the Pipeline 3 (Feature **Scaled** and **Without Tuning**).
- Pipeline 4 had **2 FP** Values lesser than Pipeline 3.
- Pipeline 4 has the **Same** AUROC Score as Pipeline 3.
- Pipeline 4 has a **Same** GINI Score as Pipeline 3.
(AUROC and GINI based on TPR/FPR, Confusion Matrices Show Small Positive/Negative Value Improvements)

About Same AUROC/GINI:

- 1) Feature Scaling Bridged/Normalized “features with larger numerical values” **[7]** giving GridSearch Optimizer Similar Transformed Datasets to Work With; and/or
- 2) Pipelines’ Relatively Small **TP** and **FP** Hard to Notice.

However, despite the Same AUROC/GINI Scores (I.e. Dependent on TPR/FPR Values), their Derived **TP/FP/TN/FN Values** can be Taken from the Confusion Matrix.

In Summary

- This Capstone was able to Fulfil the Company's need for :
 - An **Interpretable** Scorecard
 - Complex Enough to be Trained via the Logistic Regression (LR) Model with Grid Search Optimization to **Maximize the Predicted Credit Scores'** Accuracy

Deliverables:

- The **Scorecard** and **Trained Model's Pipeline** was then **Conformed** to the FICO™ Credit Scoring Range (300 to 850), **allowing the Company to Analyze** Further if the Predicted Scores are still Relevant/Usable (I.e. by being within the FICO™ Range)
- The Capstone's Workflow Design included Initial Data Pre-Processing of **145 Features** which includes:
 - **Removing** Empty Features
 - **Preparing the Model Labels** for Each Record's Default Outcome
 - **Ensuring** High Variance (reduced biases) through Dataset Splitting
 - **Transforming** the 'Train/Test Datasets' Data into Numerical Values for Compatibility with the Machine Learning Model later
 - **Feature Selection** (i.e. ANOVA F-Test and Chi-Squared) to Determine the Most Significant Features within the Dataset
 - Top 4 **Categorical** and 20 **Continuous Features** were Selected for **WOE Binning**
 - One-Hot Encoded **Eliminated Relational Value Biases** within Each Feature
 - Dataset was then **Analyzed** based on its Weight of Evidence (WOE) and Information Values prior to being WOE Binned, respectively.

In Summary

Deliverables:

- The Capstone's **Main Aim in Resolving the Gaps**, such as How Model Tuning and Feature Scaling could **Improve ML Model Performance**, was Identified in Previous Credit Scoring Studies [\[3\]](#), [\[4\]](#), [\[5\]](#).
 - These Resolutions were accomplished by Employing Grid Search Optimization (i.e. **Model Tuning**) on the LR Model with the Inclusion of Standard Scaler (i.e. **Feature Scaling**).
- **Accuracy** of the four differing ML Experimental Pipeline Configurations was **Evaluated** through Metrics like:
 - ROC Curves
 - AUROC
 - GINI Coefficient
 - TPR/FPR Confusion Matrixes
- **Results**
 - Pipeline with the Highest Accuracy was the **combined use of Grid Search and Standard Scaler**, together with the standard WOE Binning
 - This Pipeline Attained the **Most True-Positives and True-Negatives**, with the **Lowest False-Negatives and False-Positives**.
- **Creation of a Credit Scorecard** with the Chosen Pipeline producing Credit Scores based on its Customers' Loan Application Data
- **Future Improvement**
 - Potential use of the Model Tuning **Bayesian Search Optimizer** (with ROC evaluation for TPR/FPR Values)
 - Could Replace Grid Search and **Reduce Overall Time Complexity and Computation Overhead** while still providing Parameters with Highest Model Accuracy.

Objective 2.3

Score-carding with Machine Learning

Sample of Scorecard

Out[38]:

	index	Feature name	Coefficients	Original feature name	Score - Calculation	Score - Preliminary
0	0	Intercept	3.928058	Intercept	599.362503	599.0
1	1	grade:A	0.393112	grade	25.982462	26.0
2	2	grade:B	0.340824	grade	22.526535	23.0
3	3	grade:C	0.232294	grade	15.353307	15.0
4	4	grade:D	0.168589	grade	11.142754	11.0
...
103	11	tot_hi_cred_lim:>499999.95	0.000000	tot_hi_cred_lim	0.000000	0.0
104	12	mths_since_last_credit_pull_d:>56.3	0.000000	mths_since_last_credit_pull_d	0.000000	0.0
105	13	mths_since_issue_d:>93.2	0.000000	mths_since_issue_d	0.000000	0.0
106	14	mths_since_recent_inq:>18.75	0.000000	mths_since_recent_inq	0.000000	0.0
107	15	grade:E_F_G	0.000000	grade	0.000000	0.0

- Based on Coefficients of LR Model
- Fitted to FICO's Credit Scoring Range (300 to 850)

Creating a Scorecard

	Feature name	Coefficients
0	Intercept	3.928058
1	grade:A	0.393112
2	grade:B	0.340824
3	grade:C	0.232294
4	grade:D	0.168589
5	home_ownership:OTHER_NONE_RENT	0.006512
6	home_ownership:OWN	0.015960
7	home_ownership:MORTGAGE_ANY	-0.016434
8	int_rate:missing	0.000000
9	int_rate:<6.594	0.510612
10	int_rate:6.594-7.878	0.464779

- Load WOE Binned Feature Names and Coefficients from Logistic Regression Model into DataFrame

Creating a Scorecard

```
# create a list of all the reference categories, i.e. one category from each of the global features
ref_categories = ['int_rate:>20.718', 'last_pymnt_amnt:>12657.615', 'total_pymnt:>28483.595', 'acc_open_past_24mths:>9.6',
                  'inq_last_6mths:>1.6', 'num_tl_op_past_12m:>4.8', 'bc_open_to_buy:>35557.0', 'total_bc_limit:>55275.0',
                  'bc_util:>84.9', 'tot_hi_cred_lim:>499999.95', 'mths_since_last_credit_pull_d:>56.3',
                  'mths_since_issue_d:>93.2', 'mths_since_recent_inq:>18.75']
```

Adapted from WOE Binning Python Method

- Reference Column (Outliers- Max Values Outside Current Range) to be Dropped during WOE Binning

Creating a Scorecard

index		Feature name	Coefficients
0	0	Intercept	3.928058
1	1	grade:A	0.393112
2	2	grade:B	0.340824
3	3	grade:C	0.232294
4	4	grade:D	0.168589
...
103	11	tot_hi_cred_lim:>499999.95	0.000000
104	12	mths_since_last_credit_pull_d:>56.3	0.000000
105	13	mths_since_issue_d:>93.2	0.000000
106	14	mths_since_recent_inq:>18.75	0.000000
107	15	grade:E_F_G	0.000000
108 rows × 3 columns			

- Contain Indexes, Coefficients and Binned Feature Names of the **WOE Binned Features** and *optional* **Reference Categories** (have 0 Value Score Contribution).

Creating a Scorecard

Out[38]:

	index	Feature name	Coefficients	Original feature name	Score - Calculation	Score - Preliminary
0	0	Intercept	3.928058	Intercept	599.362503	599.0
1	1	grade:A	0.393112	grade	25.982462	26.0
2	2	grade:B	0.340824	grade	22.526535	23.0
3	3	grade:C	0.232294	grade	15.353307	15.0
4	4	grade:D	0.168589	grade	11.142754	11.0
...
103	11	tot_hi_cred_lim:>499999.95	0.000000	tot_hi_cred_lim	0.000000	0.0
104	12	mths_since_last_credit_pull_d:>56.3	0.000000	mths_since_last_credit_pull_d	0.000000	0.0
105	13	mths_since_issue_d:>93.2	0.000000	mths_since_issue_d	0.000000	0.0
106	14	mths_since_recent_inq:>18.75	0.000000	mths_since_recent_inq	0.000000	0.0
107	15	grade:E_F_G	0.000000	grade	0.000000	0.0

- Extract Original Feature by Splitting Feature Name by **:** symbol
- **Score – Calculation:** Calculate the Credit Score (resulting in **Continuous** Data Type)
- **Score – Preliminary:** Round Score – Calculation to **Round** Whole Numbers

Creating a Scorecard

$$\text{Score} = \text{Offset} + \text{Factor} * \ln(\text{Odds})$$

Figure 1: Score Calculation Equation of each 'Binned' Feature, adapted from Saddiq [20]

$$\text{Factor} = \frac{\text{Max FICO Score} - \text{Min FICO Score}}{\sum (\text{Max Coef of Main Feature}) - \sum (\text{Min Coef of Main Feature})}$$

Figure 2: Equation for Factor variable

- The Score of Each WOE 'Binned' Feature can be calculated by the equation in **Figure 1**.
 - The '**Odds**' Variable for Each Feature is derived from the Probability of an Attribute being "good or bad" (non-default or default) [20] i.e Coefficient of Each WOE 'Binned' Feature in Trained LR Model
 - The '**Factor**' Variable for Each 'Binned' Feature "depend[s] on the approval rate" [20] derived by the Equation in **Figure 2** Divides the Value Difference of the **Maximum and Minimum** FICO scores (300 to 850) with the **value difference of the Summation of the Maximum and Minimum Coefficients** of Each WOE Binned Feature (e.g. 'grade:A' / 'grade:B' / 'grade:C') based off the Main Feature (e.g. 'grade').
 - In **Figure 1**, the '**Offset**' variable is *optional* and comes into play when the Maximum and Minimum Scores (Multiplication of Factor and Odds) are Found to be Outside the FICO range (300 to 850).
 - When Implemented, the '**Offset**' would be **added/subtracted** from both the Multiplied Values (Maximum and Minimum Scores of the Model), resulting in a **Corrected Score** Predicted for Each 'Binned' Feature. The Customer's (FICO-based) Credit Score is Tabulated by Adding each Feature's Score together.

Creating a Scorecard

```
In [39]: # check the min and max possible scores of our scorecard
min_sum_score_prel = df_scorecard.groupby('Original feature name')['Score - Preliminary'].min().sum()
max_sum_score_prel = df_scorecard.groupby('Original feature name')['Score - Preliminary'].max().sum()
print(min_sum_score_prel)
print(max_sum_score_prel)
#Guideline from FICO: from 300 to 850

301.0
849.0
```

- Check if the Scorecard Values' Range is within FICO™ Range (300 to 850)

Creating a Scorecard

	Intercept	grade:A	grade:B	grade:C	grade:D	home_ownership:OTHER_NONE_RENT	home_ownership:OWN	home_ownership:MORTGAGE_ANY	int_rate:missing
0	1	0	0	0	0	0	0	1	0
1	1	0	0	1	0	0	0	1	0
2	1	0	0	0	0	1	0	0	0
3	1	0	1	0	0	0	0	1	0
4	1	0	1	0	0	0	1	0	0

5 rows × 92 columns

- Customers' Data Collected During the Company's Loan Application Process would then need to be Transformed with WOE Binning (in Picture) Found in the Model Pipeline.

Creating a Scorecard

```
df_scorecard['Score - Preliminary'] = df_scorecard['Score - Calculation'].round()
df_scorecard
```

Out[38]:

	index	Feature name	Coefficients	Original feature name	Score - Calculation	Score - Preliminary
0	0	Intercept	3.928058	Intercept	599.362503	599.0
1	1	grade:A	0.393112	grade	25.982462	26.0
2	2	grade:B	0.340824	grade	22.526535	23.0
3	3	grade:C	0.232294	grade	15.353307	15.0
4	4	grade:D	0.168589	grade	11.142754	11.0
...
103	11	tot_hi_cred_lim:>499999.95	0.000000	tot_hi_cred_lim	0.000000	0.0
104	12	mths_since_last_credit_pull_d:>56.3	0.000000	mths_since_last_credit_pull_d	0.000000	0.0
105	13	mths_since_issue_d:>93.2	0.000000	mths_since_issue_d	0.000000	0.0
106	14	mths_since_recent_inq:>18.75	0.000000	mths_since_recent_inq	0.000000	0.0
107	15	grade:E_F_G	0.000000	grade	0.000000	0.0

	Intercept	grade:A	grade:B	grade:C	grade:D	home_ownership:OTHER_NONE_RENT	home_ownership:OWN	home_ownership:MORTGAGE_ANY	int_rate:missing
0	1	0	0	0	0	0	0	1	0
1	1	0	0	1	0	0	0	1	0
2	1	0	0	0	0	1	0	0	0
3	1	0	1	0	0	0	0	1	0
4	1	0	1	0	0	0	1	0	0

5 rows × 92 columns

- Reshape Customers' Dataset from (... , 92) to (108, ...) to Perform **Dot Matrix Multiplication**

Creating a Scorecard

```
In [46]: # matrix dot multiplication of test set with scorecard scores  
y_scores = X_test_woe_transformed.dot(scorecard_scores)  
y_scores.head()
```

Out[46]:

	0
0	623.0
1	421.0
2	414.0
3	643.0
4	506.0

- Perform Dot Matrix Multiplication to get **Final Credit Scores**

Thank You

Q&A